

Literal Capture: Auditing Non-Generative Decision Models on Bangla Idioms

Anonymous ACL submission

Abstract

Decision models — non-autoregressive scorers that return a calibrated probability over caller-supplied options and emit no text — are increasingly deployed as routers, classifiers and guardrails. They generate nothing, so the usual explainability toolkit of rationales, chain-of-thought and token attribution does not apply, and idioms make token-level attribution ill-posed in principle: a non-compositional expression has no word whose removal is informative. We therefore audit three such models entirely through *interventions on the option set*, using 10,359 expert-annotated Bangla idioms that ship with both a gold figurative gloss and a gold literal gloss. Our instrument holds the item, the instruction and k fixed and varies only the distractor family, which isolates one failure mode at a time and needs no new annotation. Three results hold across an encoder scorer, a decoder scorer and a closed commercial API. First, *literal capture* is severe and does not decrease with capability: inserting the idiom’s own literal gloss drives accuracy far below chance, and the model that is better on every other condition is the more captured one. Second, confidence is anti-diagnostic under this shift: both open models are more confident when captured than when correct, and confidence-gated selective prediction is worse than random gating. Third, option-order sensitivity dissociates cleanly by architecture: 83.7% versus 10.7% of items change their chosen gloss across the 24 orderings of the same four options. A semantics-free control shows the strongest model reporting 0.815 mean confidence at 0.254 accuracy. We release all artifacts.

1 Introduction

A *decision model* is given a question and a list of candidate options and returns a probability distribution over exactly those options. It is not autoregressive, it does not sample, and it produces no text. This class has grown quickly as infrastructure —

routing, intent classification, content gating, tool selection — precisely because it is cheap, bounded and returns a number a caller can threshold on.

That bounded interface is also an explainability problem. There is no rationale to inspect, no chain of thought to measure for faithfulness (Jacovi and Goldberg, 2020; Turpin et al., 2023), and in the commercial case no weights at all. What is left is the behavioural surface: the caller controls the question and the option list, and sees a probability vector. Any audit of this class has to be built out of that.

The obvious fallback — attributing the decision to input tokens (Ribeiro et al., 2016; Lundberg and Lee, 2017) — is not merely unreliable here (Jain and Wallace, 2019; Adebayo et al., 2018; Hooker et al., 2019); for our domain it is ill-posed. An idiom is non-compositional by definition (Nunberg et al., 1994): the words of *kick the bucket* do not carry fractions of its meaning, so a Shapley value over them attributes a quantity that does not exist. The expression has to be intervened on as a unit, or not at all.

We take the option set as the intervention surface instead. Bangla *bāgdhārā* (idioms) are well suited to this because the expert resource we use ships, for each idiom, both a gold figurative meaning and a gold *literal* gloss — a ready-made, human-certified counterfactual distractor that no annotation effort of ours produced. Dropping that gloss into the option list turns a benchmark question into a targeted probe: does the model read the expression as an idiom, or as its words?

Contributions.

1. A meaning-choice instrument for decision models in which the *distractor family* is the independent variable and everything else is held fixed, requiring no new annotation (§3).
2. **Literal capture does not decrease with capa-**

bility. We define the Literal Attraction Rate and show a double dissociation: the model that is better on every non-literal condition is substantially more captured (§5).

3. **Confidence is anti-diagnostic under this shift,** across three different calibration philosophies, with confidence-gated selective prediction worse than random gating (§6).
4. **Option-order dependence dissociates by architecture,** cleanly and at scale, with a methodological consequence for anyone evaluating this class (§8).
5. A semantics-free control that bounds what each model reports when there is provably nothing to report (§7).

2 Background

Idioms in NLP. Idiom resources and tasks are well established (Haagsma et al., 2020; Tayyar Madabushi et al., 2022; Chakrabarty et al., 2022), and the difficulty of figurative language for distributional models is long-noted (Shwartz and Dagan, 2019). Psycholinguistics frames the phenomenon as competition between a literal and a figurative route whose resolution depends on decomposability and context (Cacciari and Tabossi, 1988; Gibbs et al., 1989; Titone and Connine, 1999). Our D-LIT intervention is a deliberate behavioural analogue of that competition: we place both routes in the option list and read off which one wins.

Evaluating option scorers. Option-order sensitivity and selection bias are documented for prompted generative LMs (Pezeshkpour and Hruschka, 2024; Wei et al., 2024; Zheng et al., 2024; Zhao et al., 2021). We are not aware of a comparison of option-marker against independent-option scoring in small non-autoregressive models over the full permutation set, which is what §8 provides.

Calibration and selective prediction. Proper scoring and calibration measurement are standard (Brier, 1950; Gneiting and Raftery, 2007; Guo et al., 2017; Nixon et al., 2019; Kumar et al., 2019), as is the observation that calibration degrades under distribution shift (Ovadia et al., 2019; Jiang et al., 2021). Selective prediction gives the operational test (Geifman and El-Yaniv, 2017, 2019; Traub et al., 2024). We report the combination that matters for deployment: accuracy *decreasing* with

confidence, and a confidence gate underperforming a random one.

Behavioural testing. Our design is in the tradition of behavioural checklists, contrast sets and counterfactual minimal pairs (Ribeiro et al., 2020; Gardner et al., 2020; Kaushik et al., 2020; Naik et al., 2018; Jia and Liang, 2017). The contribution is to carry that tradition into a model class with no text output, where it is one of the few instruments available.

Low-resource evaluation. Bangla is served by strong encoders and corpora (Bhattacharjee et al., 2022; Kakwani et al., 2020), and cross-lingual benchmarks establish the transfer gap (Hu et al., 2020; Ruder et al., 2021; Adelani et al., 2021). §11 localises our failure to the Bangla side rather than to figurativeness itself.

3 The meaning-choice instrument

Data. We use a corpus of 10,359 expert-annotated Bangla idioms. Each row carries the idiom, a gold figurative meaning in Bangla, an expert literal gloss, semantic tags, alternative surface forms certified as meaning-preserving, and a gold carrier sentence. We add no annotation; every intervention below is assembled from fields that already exist.

Task. Given an idiom, choose its gold figurative meaning from $k = 4$ candidates. Chance is 0.250 throughout. The **distractor family is the independent variable**; the item, the instruction, k , and per-item option-order randomisation (seeded by item id) are held fixed across families, so a paired test across families is a test on byte-identical questions.

Distractor families. D-RAND draws other idioms’ figurative glosses at random. D-SURF draws glosses that share surface tokens with the gold. D-SEM draws glosses from idioms sharing semantic tags. D-LIT replaces one distractor with the idiom’s own expert literal gloss. D-HARD combines all three pressures.

Literal Attraction Rate. For the families containing a literal gloss, let ℓ_i be the literal option and g_i the gold:

$$\text{LAR} = \frac{1}{N} \sum_{i=1}^N \mathbf{1} \left[\arg \max_o p_i(o) = \ell_i \right].$$

LAR is not 1 – accuracy: it isolates the share of errors that are the specific error of reading the idiom

literally, and it is bounded below by nothing — a model that never reads literally scores 0 while still being wrong.

Seven interventions. E1 is the distractor ladder above. E2 swaps the idiom for a certified synonymous surface form. E3 perturbs the idiom as a unit (shuffle, reverse, delete one word, substitute one word, embed in its carrier sentence) with the options held fixed. E4 swaps the task language. E5 measures calibration and selective prediction. E6 ablates the option text itself (truncate, scramble, translate, strip to bare labels). E7 runs all $4! = 24$ orderings of a fixed option set.

Statistics. All intervals are 95% percentile bootstrap over items (2,000 resamples). Paired comparisons use McNemar’s exact test. ECE uses 15 equal-width bins; we also report AURC against a random-gate baseline.

Pairing verification. The two open arms were run in separate accounts from independent notebooks, each reconstructing the corpus and the task instrument from a hardcoded contract. On all 4,440 overlapping rows the gold key, gold position, idiom string and literal-distractor key agree at 1.000. Every paired test below is therefore on identical questions, not on re-derivations that happen to look alike.

4 Models

We audit three decision models spanning the design space (Table 1). **Laya-ml** is a bidirectional encoder with an option-marker scorer: all options are encoded together and scored from marker positions. **Lod-lille** is a decoder backbone with independent-option scoring, where each option attends only to the shared state, the question and itself from the same start position — order-independent by construction in `fp32`. **Jev-1.13** is a closed commercial model behind a decisions API; we know its interface and nothing else. It is not multilingual by documentation, which we exploit rather than regret: it provides the capability contrast that makes §5 a dissociation rather than an anecdote.

5 Literal capture does not decrease with capability

Table 2 and Figure 1 give the ladder for the two open arms on the 888 shared items. Lod-lille is the stronger model on every non-literal condition, by

	Laya-ml	Lod-lille	Jev-1.13
backbone	mmBERT	Qwen3-0.6B	closed
scoring	marker	independent	unknown
params	322M	0.6B	—
items	8,876	888	8,761
latency	27.2 ms	86.1 ms	API
weights	open	open	closed

Table 1: The three arms. Cell tint identifies the model throughout the paper: cyan Laya-ml, magenta Lod-lille, yellow Jev-1.13. Lod-lille ran a fixed 10% random subset; Appendix B bounds the resulting drift at 0.016.

Distractors	Laya	Lod	Δ
random	0.367	0.484	+0.133***
surface	0.338	0.412	+0.073***
tag	0.333	0.446	+0.120***
+ literal	0.151	0.066	−0.070***
hard	0.144	0.063	−0.086***
LAR (lit.)	0.644	0.885	
LAR (hard)	0.631	0.855	
mean p (lit.)	0.464	0.726	
mean p (gold)	0.195	0.128	

Table 2: The distractor ladder and literal attraction, McNemar on the 888 shared items (***) $p < 0.001$. Bold marks the better of the two arms per row. The sign of Δ flips at the horizontal rule: capability on the surrounding task does not protect against literal capture, and here is associated with more of it.

up to +0.133. Insert the idiom’s own literal gloss and the ordering *reverses*: Lod-lille falls to 0.066, less than half of Laya-ml’s already-below-chance 0.151, and barely a quarter of chance. The same reversal holds for D-HARD.

This is the central result. Capability on the surrounding task does not protect against literal capture. LAR therefore measures something accuracy does not: a model can be selected as “better” on a conventional benchmark while being strictly worse at the thing that benchmark is a proxy for. The error is not diffuse either — the full decomposition under D-LIT is gold 0.151 / literal 0.644 / other 0.207 for Laya-ml and 0.066 / 0.885 / 0.048 for Lod-lille. On the clearest items Laya-ml places $p = 1.000$ on the literal gloss and 0.000 on the gold.

A confound we do not hide. Jev-1.13 reaches $LAR = 0.047$, two orders of magnitude below the open arms. That number is *English*. The shared literal pool exists only for the 3,265 items with an English literal gloss, and the commercial arm ran it in English while the open arms ran it in Bangla. Figure 2 draws the full language $\times$ model matrix




Figure 1: The distractor ladder (left) and literal attraction rate (right) for the two open arms. **Cyan** is Laya-ml ($n = 8,876$), **magenta** is Lod-lille ($n = 888$); the black dashed line is chance (0.250) and error bars are 95% bootstrap CIs. Surface- and tag-matched distractors cost a few points. Inserting the idiom’s own literal gloss (D-LIT) collapses both models far below chance, and collapses the *stronger* model harder: Lod-lille leads Laya-ml by +0.133 on random distractors and trails it by -0.070 once the literal gloss is present. On the right, both models prefer the literal reading to the figurative one on the majority of items.




Figure 2: Literal attraction on the shared 3,265-item pool, by model and task language. Solid bars are Bangla, the hatched **yellow** bar is Jev-1.13’s *English* condition, and the light grey hatched blocks running past the top of the axis mark cells that were **never run** — they are not values of 1.0. Jev-1.13’s low attraction cannot be read as a Bangla result; the experiment that would resolve the confound is exactly the two missing cells.

Model	n	Accuracy	95% CI
Laya-ml	8,876	0.367	[.357, .377]
Lod-lille	888	0.484	[.453, .518]
Jev-1.13	8,761	0.736	[.727, .745]

Table 3: The only condition identical across all three arms: Bangla idiom $\rightarrow$ Bangla gloss, random distractors. All pairwise McNemar tests on the overlapping items are significant: Jev–Laya +0.369, Jev–Lod +0.259, Lod–Laya +0.133 (all $p < 10^{-9}$).

including the two cells that were never run, so the reader can see the shape of the missing experiment rather than infer a clean three-way comparison that we did not perform. A second caveat belongs with it: the pool is not neutral. Laya-ml’s in-pool LAR is 0.747 against 0.583 out of pool, while its D-RAND accuracy differs by only 0.008 — items that happen to have an English literal gloss are the more literally-attracting ones.

The one matched three-way comparison. Exactly one condition is identical across all three arms: Bangla idiom, Bangla gloss, random distractors.

Figure 3 and Table 3 report it. The commercial model is far ahead (0.736 against 0.367 and 0.484), which is the honest capability ordering, and the one against which the calibration results of §6 and §7 should be read.

6 Confidence is anti-diagnostic

Table 4 gives the mean probability assigned to the chosen option, split by outcome, on the literal-bearing families. Both open models are *more confident when captured than when correct*, Lod-lille extremely so (0.787 against 0.453). The reliability curves in Figure 4 slope *downward* under D-LIT in both.

Selective prediction therefore inverts. Restricting Laya-ml’s D-LIT to the most confident 5% of items gives 0.007 accuracy against 0.151 at full coverage, and the confidence-gated AURC is worse than a random gate for both open arms. The failure is not an artifact of one calibration philosophy: Laya-ml ships fitted per-bucket temperatures,




Figure 3: The matched three-way comparison. Bars are, left to right, **cyan** Laya-ml, **magenta** Lod-lille, **yellow** Jev-1.13; black dashed line is chance, error bars 95% bootstrap CIs. This capability ordering is the baseline against which the calibration failures in §6–§7 should be read.

Outcome	Laya	Lod	Jev
correct	0.432	0.453	0.935
captured	0.596	0.787	0.737
other error	0.421	0.399	0.746
ECE (D-RAND)	0.071	0.073	0.005
ECE (D-LIT)	0.385	0.679	0.014
AURC (conf.)	0.924	0.989	0.013
AURC (random)	0.849	0.933	0.069

Table 4: Confidence by outcome on the literal-bearing families, and calibration. Bold marks, per model, whether the model is more confident when correct or when captured. Both open arms are more confident when captured; their confidence-gated AURC is *worse* than the random-gate baseline in the grey row, which inverts selective prediction. Jev-1.13’s English condition does not show the inversion.

Lod-lille deliberately ships none, and Lod-lille further ships a *separate confidence head* whose documented purpose is estimating whether the answer is in the option list at all. The head does not rescue it — mean head confidence is 0.631 on captured items against 0.295 on correct ones, the same inversion.

Jev-1.13 is the interesting exception, and it is only an exception in one direction. Its reliability is near-ideal in English (ECE 0.005–0.047) and its confidence-gated AURC beats random by an order of magnitude. That makes §7 harder to dismiss, not easier.

7 A semantics-free control

The O4 ablation replaces every option’s text with a bare label. No semantic signal remains, so accuracy is pinned at chance by construction. Whatever confidence survives is therefore pure overconfidence,

	Laya	Lod	Jev
answer changes	0.837	0.107	0.133
correctness varies	0.686	0.070	0.120
distinct glosses	2.83	1.14	1.15
perfectly stable	0.163	0.893	0.867
$p(\text{gold})$ range	0.346	0.029	0.128
selection spread	0.226	0.052	0.072

Table 5: Option-order sensitivity over all 24 orderings of the same four options ($n = 992 / 299 / 300$ items). Bold marks the worst arm per row. The marker-based scorer changes its chosen gloss on 83.7% of items; the independent-option scorer, which is order-independent by construction in fp32, changes it on 10.7% and very nearly delivers what its design claims.

with no modelling assumption required to interpret it. Figure 5 reports it.

The ordering reverses against Table 4. The best-calibrated model under real options is the worst when the options carry nothing: Jev-1.13 reports 0.815 mean confidence at 0.254 accuracy, an ECE of 0.561. Lod-lille, the weakest of the three on accuracy, is the only one whose reported confidence collapses appropriately (0.282). We read this as the operational point of the paper: a calibration number measured on in-distribution options says very little about what the model will report when the option list is degenerate, and a production caller thresholding on a returned probability is exposed to exactly that case.

A useful side-check comes with it. McNemar between the two open arms under O4 is not significant ($p = 0.76$); the instrument’s null behaves identically in both, which is evidence that the task is measuring option semantics rather than an artifact of either scorer.

8 Option order is architecture-specific

We run every item under all $4! = 24$ orderings of a fixed option set and record the identity of the chosen gloss, recovered as $\text{perm}[\text{slot}(\hat{y})]$.¹

Table 5 and Figure 6 give a clean dissociation: a $7.8\times$ difference in answer stability and a $12\times$ difference in probability stability, tracking the scoring architecture rather than model size, backbone or capability. Lod-lille’s design claim — each option attends only to the shared state, its question and

¹An earlier version of this analysis counted distinct option keys (slot letters) and reported near-total instability for Laya-ml. That metric is an artifact: a key necessarily changes when the same gloss moves slots. All numbers here use the corrected gloss-identity metric; Laya-ml’s figure drops from 1.000 to 0.837 and the finding stands.




Figure 4: Reliability, 15 equal-width bins, panels left to right **Laya-ml**, **Lod-lille**, **Jev-1.13**. Dark grey circles are the random-distractor condition; model-coloured squares are the +literal condition. The grey dashed diagonal is perfect calibration. Under D-LIT the two open models’ curves bend *downward* — higher stated confidence buys lower accuracy — while Jev-1.13, running in English, stays on the diagonal. Reading the left two panels: the model is most wrong exactly where it reports being most sure.




Figure 5: The semantics-free control. Per model, the **grey** bar is accuracy (with 95% CI) and the **coloured** bar is mean reported confidence; the black dashed line is chance. Because the options carry no information, accuracy is at chance by construction and the gap between the two bars is pure overconfidence: 0.250 for Laya-ml, 0.049 for Lod-lille, and 0.561 for Jev-1.13 — which reports 0.815 confidence while answering at 0.254. Jev-1.13’s condition is English.

itself from the same start position — is very nearly borne out. Laya-ml’s option-marker scorer is not, confirming an upstream issue report at scale.

The residual positional bias follows (Figure 7). With the gold placed uniformly, Laya-ml selects slot B 37.2% of the time against 14.6% for slot A, and its accuracy is 0.152 when the gold sits in A against 0.403 in B — a $2.6\times$ swing from position alone.

Methodological consequence. Any evaluation of this model class that fixes option order is partly measuring position. All results in this paper randomise order per item, seeded by item id, and are

therefore averages over position rather than measurements at one of them. We suggest this should be a default for the class, not a robustness check.

9 Options are read as a bag of words

The O-series ablates the option text while holding the idiom fixed (Figure 8). Two readings matter. First, truncating each option to six words and *scrambling the words inside it* cost almost nothing: Laya-ml retains 0.967 and 0.925 of its signal, Lod-lille 0.909 and 0.948, Jev-1.13 is flat to within noise. Whatever all three models are doing on the option side is close to order-free lexical matching. Second, the drop to the grey semantics-free bar is the only large one in the figure, which is what certifies the instrument: the task really is measuring option semantics, it is just measuring them as a bag.

10 Neither open model treats the idiom as a unit

E3 perturbs the idiom while holding the option list fixed (Figure 9). This is the experiment that replaces token attribution: since an idiom has no word-level meaning to attribute to, we instead ask whether the model behaves as though it stores the expression as a unit. A model that did would collapse when the unit is broken.

Neither does. Word-shuffling an idiom leaves Laya-ml’s accuracy statistically unchanged (0.158 $\rightarrow$ 0.168) and reversal actually helps slightly (0.180); Lod-lille moves within noise on its 180 items. Answer agreement with the intact form




Figure 6: Order sensitivity. *Left*: how many distinct glosses each model selects across the 24 orderings of one option set; within each group bars run **cyan** Laya-ml, **magenta** Lod-lille, **yellow** Jev-1.13, and a value of 1 means perfect stability. *Right*: per model, the **coloured** bar is the share of items whose chosen gloss is unstable and the **grey** bar is the mean within-item range of $p(\text{gold})$. Laya-ml is unstable on 83.7% of items with a mean $p(\text{gold})$ swing of 0.346; the other two are near-stable.




Figure 7: Share of selections falling on each option slot, panels left to right **Laya-ml**, **Lod-lille**, **Jev-1.13**. The gold option is placed uniformly at random, so the black dashed line at 0.25 is the expectation under no positional bias and any departure from it is bias. Laya-ml takes slot B 37.2% of the time and slot A 14.6%; the other two show only a mild primacy gradient.

under shuffling is 0.742 for Laya-ml and 0.917 for Lod-lille — Lod-lille is more *stable*, consistent with §8, while being no more *compositional*.

The one manipulation that clearly helps both is embedding the idiom in its gold carrier sentence: +0.080 and +0.117. Context, not the idiom string, carries the usable signal — which is itself the diagnosis. These models are not retrieving a stored figurative entry; they are matching whatever lexical material is available, and a carrier sentence simply supplies more of it.

11 The failure localises to Bangla

Holding the model fixed and swapping only the task language (Figure 10) moves every arm by a large margin: Laya-ml +0.246, Lod-lille +0.380, Jev-1.13 +0.233 (all McNemar $p < 0.001$). Lod-lille

reaches 0.864 in English against 0.484 in Bangla, so the figurative concepts are largely present — it simply cannot reach them through Bangla. Its documentation states that training evaluation was English-only, and this is what that looks like from outside the model.

Laya-ml’s control arm closes the argument: its English-only checkpoint run on Bengali script sits at 0.283, barely above chance, which reproduces a documented failure mode and rules out the possibility that the Bangla result is an artifact of the scorer rather than of the language.

12 Discussion

Three of our four claims are properties of the *class*, established by replication across architecturally unrelated models rather than by a single model. Lit-




Figure 8: Option-side ablation, panels left to right **Laya-ml**, **Lod-lille**, **Jev-1.13**. In each panel the fifth, **grey** bar is the semantics-free control of §7; the whole Jev-1.13 panel is hatched because its task is in **English**. Error bars are 95% bootstrap CIs and the black dashed line is chance. Truncating and word-scrambling the option text cost almost nothing in any of the three models.

eral capture, anti-diagnostic confidence and the no-signal control all reproduce across an encoder scorer and a decoder scorer, and the commercial arm — which escapes the first two in English — fails the third most severely of the three. Option-order dependence is the exception, and informatively so: it tracks the scoring architecture, and the architecture that claims order-independence very nearly delivers it.

For practitioners the implications are concrete. A returned probability from this model class should not be thresholded without a measurement of what the model returns on a degenerate option list, because that number can be high while accuracy is at chance. A benchmark score should not be read as evidence against literal capture, because the two move in opposite directions here. And an evaluation that fixes option order is partly reporting position.

For the low-resource story, §11 is the sharpest form of a familiar result: these models hold the concepts and cannot reach them through Bangla. That is a different problem from not having the concepts, and it suggests a different remedy.

Limitations

All three models are evaluated zero-shot; nothing here speaks to a fine-tuned model of any of these families, and the open arms’ documentation is explicit that base checkpoints are weak on typed-decision benchmarks.

Laya-ml’s run had its per-bucket temperatures unset, so its ECE figures are raw. Lod-lille ships no temperature by design, so both open arms are effectively untempered and comparable on that axis; Jev-1.13’s internal calibration is unknown to us.

Lod-lille ran a fixed 10% subset (888 items), so its intervals are roughly ± 3 points and its stratified cells are thin. Appendix B bounds the subset bias at 0.016 and shows it changes no sign or ordering. E3 on the Lod-lille arm is 180 items and several of its perturbation contrasts are not individually significant. E7 covers 992 / 299 / 300 items rather than the full sets.

The cross-model LAR comparison in §5 is confounded with language, as Figure 2 states explicitly; the two missing cells are the experiment that would resolve it and we did not run them. The shared literal pool is also not a neutral sample of the corpus.

Distractors are drawn from the same corpus, so difficulty is corpus-relative rather than absolute. Jev-1.13 is a commercial endpoint and may change without notice; our results are a snapshot, and we report the exact call count and cost in Appendix D so the snapshot is at least auditable.

Finally, no figure in this paper contains Bangla glyphs: the plotting toolchain available to us could not shape Bengali script. All categories are named in Latin script and qualitative Bangla examples are confined to the text.




Figure 9: Unit integrity: the idiom is perturbed, the options are held fixed. **Cyan** is Laya-ml, **magenta** is Lod-lille; error bars 95% bootstrap CIs, black dashed line chance (left panel). *Left*: accuracy. *Right*: agreement with the answer given on the intact idiom. Scrambling and reversing the word order of a non-compositional expression does not meaningfully hurt either model, which is the bag-of-words signature. Only the carrier-sentence variant clearly helps.

Ethics Statement

The corpus is a published, expert-annotated resource of idioms; it contains no personal data and no content about identifiable individuals. Our interventions rewrite option lists and perturb idiom strings, and produce no new annotation and no generated text.

The work audits two openly released models and one commercial API. We report failure modes in all three, including the vendor’s. We have taken care to separate what is measured from what is inferred — in particular, Figure 2 marks the cells we did not run rather than letting a favourable contrast stand unqualified, and §5 states the pool selection effect against our own cleanest number. Commercial endpoints are mutable; readers should treat the Jev-1.13 results as dated rather than as a standing property of the product.

Bangla is under-served relative to its speaker population, and we think an audit that locates a failure in the language rather than in the task is the more useful kind to publish. We release the instrument, all per-item outputs and the analysis code so the measurements can be repeated on other models and other languages.

References

Julius Adebayo, Justin Gilmer, Michael Muelly, Ian Goodfellow, Moritz Hardt, and Been Kim. 2018. [Sanity checks for saliency maps](#). In *Advances in Neural Information Processing Systems 31*.

David Ifeoluwa Adelani, Jade Abbott, Graham Neubig, et al. 2021. [Masakhaner: Named entity recognition for african languages](#). *Transactions of the Association for Computational Linguistics*.

Abhik Bhattacharjee, Tahmid Hasan, Kazi Samin, Md Saiful Islam, M. Sohel Rahman, Anindya Iqbal, and Rifat Shahriyar. 2022. [Banglabert: Language model pretraining and benchmarks for low-resource language understanding evaluation in bangla](#). In *Findings of NAACL 2022*.

Glenn W. Brier. 1950. [Verification of forecasts expressed in terms of probability](#). *Monthly Weather Review*.

Cristina Cacciari and Patrizia Tabossi. 1988. [The comprehension of idioms](#). *Journal of Memory and Language*.

Tuhin Chakrabarty, Arkadiy Saakyan, Debanjan Ghosh, and Smaranda Muresan. 2022. [Flute: Figurative language understanding through textual explanations](#). In *Proceedings of EMNLP 2022*.

Matt Gardner, Yoav Artzi, Victoria Basmova, Jonathan Berant, et al. 2020. [Evaluating models’ local decision boundaries via contrast sets](#). In *Findings of EMNLP 2020*.

Yonatan Geifman and Ran El-Yaniv. 2017. [Selective classification for deep neural networks](#). In *Advances in Neural Information Processing Systems 30*.

Yonatan Geifman and Ran El-Yaniv. 2019. [Selectivenet: A deep neural network with an integrated reject option](#). In *Proceedings of ICML 2019*.

Raymond W. Gibbs, Nandini P. Nayak, and Cooper Cutting. 1989. [How to kick the bucket and not decompose: Analyzability and idiom processing](#). *Journal of Memory and Language*.




Figure 10: Accuracy by task language, grouped by model: **cyan** Laya-ml (four bars), **magenta** Lod-lille, **yellow** Jev-1.13. **Solid** bars are Bangla input, **hatched** bars are English input; Laya-ml’s third and fourth bars use the English-only checkpoint, the fourth being the control. Error bars 95% bootstrap CIs, black dashed line chance. Every model gains 20–45 points from presenting the same idiom in English, and the English checkpoint on Bengali script (fourth bar) sits near chance.

Tilman Gneiting and Adrian E. Raftery. 2007. [Strictly proper scoring rules, prediction, and estimation](#). *Journal of the American Statistical Association*.

Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. 2017. [On calibration of modern neural networks](#). In *Proceedings of ICML 2017*.

Hessel Haagsma, Johan Bos, and Malvina Nissim. 2020. [Magpie: A large corpus of potentially idiomatic expressions](#). In *Proceedings of the Twelfth Language Resources and Evaluation Conference*.

Sara Hooker, Dumitru Erhan, Pieter-Jan Kindermans, and Been Kim. 2019. [A benchmark for interpretability methods in deep neural networks](#). In *Advances in Neural Information Processing Systems 32*.

Junjie Hu, Sebastian Ruder, Aditya Siddhant, Graham Neubig, Orhan Firat, and Melvin Johnson. 2020. [Xtreme: A massively multilingual multi-task benchmark for evaluating cross-lingual generalization](#). In *Proceedings of ICML 2020*.

Alon Jacovi and Yoav Goldberg. 2020. [Towards faithfully interpretable nlp systems: How should we define and evaluate faithfulness?](#) In *Proceedings of ACL 2020*.

Sarthak Jain and Byron C. Wallace. 2019. [Attention is not explanation](#). In *Proceedings of NAACL 2019*.

Robin Jia and Percy Liang. 2017. [Adversarial examples for evaluating reading comprehension systems](#). In *Proceedings of EMNLP 2017*.

Zhengbao Jiang, Jun Araki, Haibo Ding, and Graham Neubig. 2021. [How can we know when language models know? on the calibration of language models for question answering](#). *Transactions of the Association for Computational Linguistics*.

Divyanshu Kakwani, Anoop Kunchukuttan, Satish Golla, N. C. Gokul, Avik Bhattacharyya, Mitesh M. Khapra, and Pratyush Kumar. 2020. [IndicNLPsuite: Monolingual corpora, evaluation benchmarks and pre-trained multilingual language models for indian languages](#). In *Findings of EMNLP 2020*.

Divyansh Kaushik, Eduard Hovy, and Zachary C. Lipton. 2020. [Learning the difference that makes a difference with counterfactually-augmented data](#). In *International Conference on Learning Representations*.

Ananya Kumar, Percy S. Liang, and Tengyu Ma. 2019. [Verified uncertainty calibration](#). In *Advances in Neural Information Processing Systems 32*.

Scott M. Lundberg and Su-In Lee. 2017. [A unified approach to interpreting model predictions](#). In *Advances in Neural Information Processing Systems 30*.

Aakanksha Naik, Abhilasha Ravichander, Norman Sadeh, Carolyn Rose, and Graham Neubig. 2018. [Stress test evaluation for natural language inference](#). In *Proceedings of COLING 2018*.

Jeremy Nixon, Michael W. Dusenberry, Linchuan Zhang, Ghassen Jerfel, and Dustin Tran. 2019. [Measuring calibration in deep learning](#). In *CVPR Workshops 2019*.

Geoffrey Nunberg, Ivan A. Sag, and Thomas Wasow. 1994. [Idioms](#). *Language*.

Yaniv Ovadia, Emily Fertig, Jie Ren, Zachary Nado, D. Sculley, Sebastian Nowozin, Joshua Dillon, Balaji Lakshminarayanan, and Jasper Snoek. 2019. [Can you trust your model’s uncertainty? evaluating predictive uncertainty under dataset shift](#). In *Advances in Neural Information Processing Systems 32*.

Pouya Pezeshkpour and Estevam Hruschka. 2024. Large language models sensitivity to the order of options in multiple-choice questions. In *Findings of NAACL 2024*.

Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. 2016. “why should i trust you?”: Explaining the predictions of any classifier. In *Proceedings of KDD 2016*.

Marco Tulio Ribeiro, Tongshuang Wu, Carlos Guestrin, and Sameer Singh. 2020. Beyond accuracy: Behavioral testing of nlp models with checklist. In *Proceedings of ACL 2020*.

Sebastian Ruder, Noah Constant, Jan Botha, Aditya Siddhant, Orhan Firat, et al. 2021. Xtreme-r: Towards more challenging and nuanced multilingual evaluation. In *Proceedings of EMNLP 2021*.

Vered Shwartz and Ido Dagan. 2019. Still a pain in the neck: Evaluating text representations on lexical composition. *Transactions of the Association for Computational Linguistics*.

Harish Tayyar Madabushi, Edward Gow-Smith, Carolina Scarton, and Aline Villavicencio. 2022. Semeval-2022 task 2: Multilingual idiomaticity detection and sentence embedding. In *Proceedings of SemEval 2022*.

Debra A. Titone and Cynthia M. Connine. 1999. On the compositional and noncompositional nature of idiomatic expressions. *Journal of Pragmatics*.

Julian Traub et al. 2024. Overcoming common flaws in the evaluation of selective classification systems. In *Advances in Neural Information Processing Systems 37*.

Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. 2023. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems 36*.

Sheng-Lun Wei, Cheng-Kuang Wu, and Hsin-Hsi Chen. 2024. Unveiling selection biases: Exploring order and token sensitivity in large language models. In *Findings of ACL 2024*.

Zihao Zhao, Eric Wallace, Shi Feng, Dan Klein, and Sameer Singh. 2021. Calibrate before use: Improving few-shot performance of language models. In *Proceedings of ICML 2021*.

Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. 2024. Large language models are not robust multiple choice selectors. In *International Conference on Learning Representations*.

A Surface invariance

The corpus ships alternative surface forms certified by annotators as meaning-preserving. A model

reading meaning should be invariant across such a pair. Laya-ml agrees with itself on 0.657 of 3,143 pairs (2,423 items) and Lod-lille on 0.677 of 310 pairs (234 items): about a third of the time, a meaning-preserving rewrite changes the answer. Mean Jensen–Shannon divergence between the two probability vectors is 0.039 and 0.085.

B Subset representativeness

Lod-lille ran a fixed 10% random subset. Table 6 compares Laya-ml’s accuracy on those same 888 items against its own full 8,876.

Family	full	subset	drift
D-RAND	0.367	0.351	−0.016
D-SURF	0.338	0.339	+0.001
D-SEM	0.333	0.325	−0.007
D-LIT	0.151	0.136	−0.015
D-HARD	0.144	0.149	+0.004

Table 6: Laya-ml on the 10% subset against its own full set. Maximum drift 0.016. Two of five subset estimates fall marginally outside the full-set bootstrap CI, which is expected at $n = 888$ and changes no sign or ordering. Cite the full-set figures as the precise ones.

C Full option ablation

Table 7 gives the numbers behind Figure 8.

Option text	Laya	Lod	Jev
full gloss	0.365	0.492	0.970
truncated	0.361	0.468	0.958
scrambled	0.357	0.479	0.974
other lang.	0.347	0.398	0.963
no semantics	0.256	0.233	0.254

Table 7: Option-side ablation. The grey row is the semantics-free control, at chance by construction in all three arms. Jev-1.13’s column is English.

D Reproducibility

The two open arms ran on a single T4 GPU, Laya-ml in `bf16` at 27.2 ms per call over 8,876 items and Lod-lille in `fp32` at 86.1 ms per call over a fixed 888-item subset. The commercial arm consumed 56,533 API calls at a total cost of \$0.93. All intervals are 95% percentile bootstrap over items with 2,000 resamples at a fixed seed; ECE uses 15 equal-width bins; paired tests are McNemar’s exact test. Per-item outputs, the analysis code, the figure sources and the exact plotted values for every figure are released with the paper.